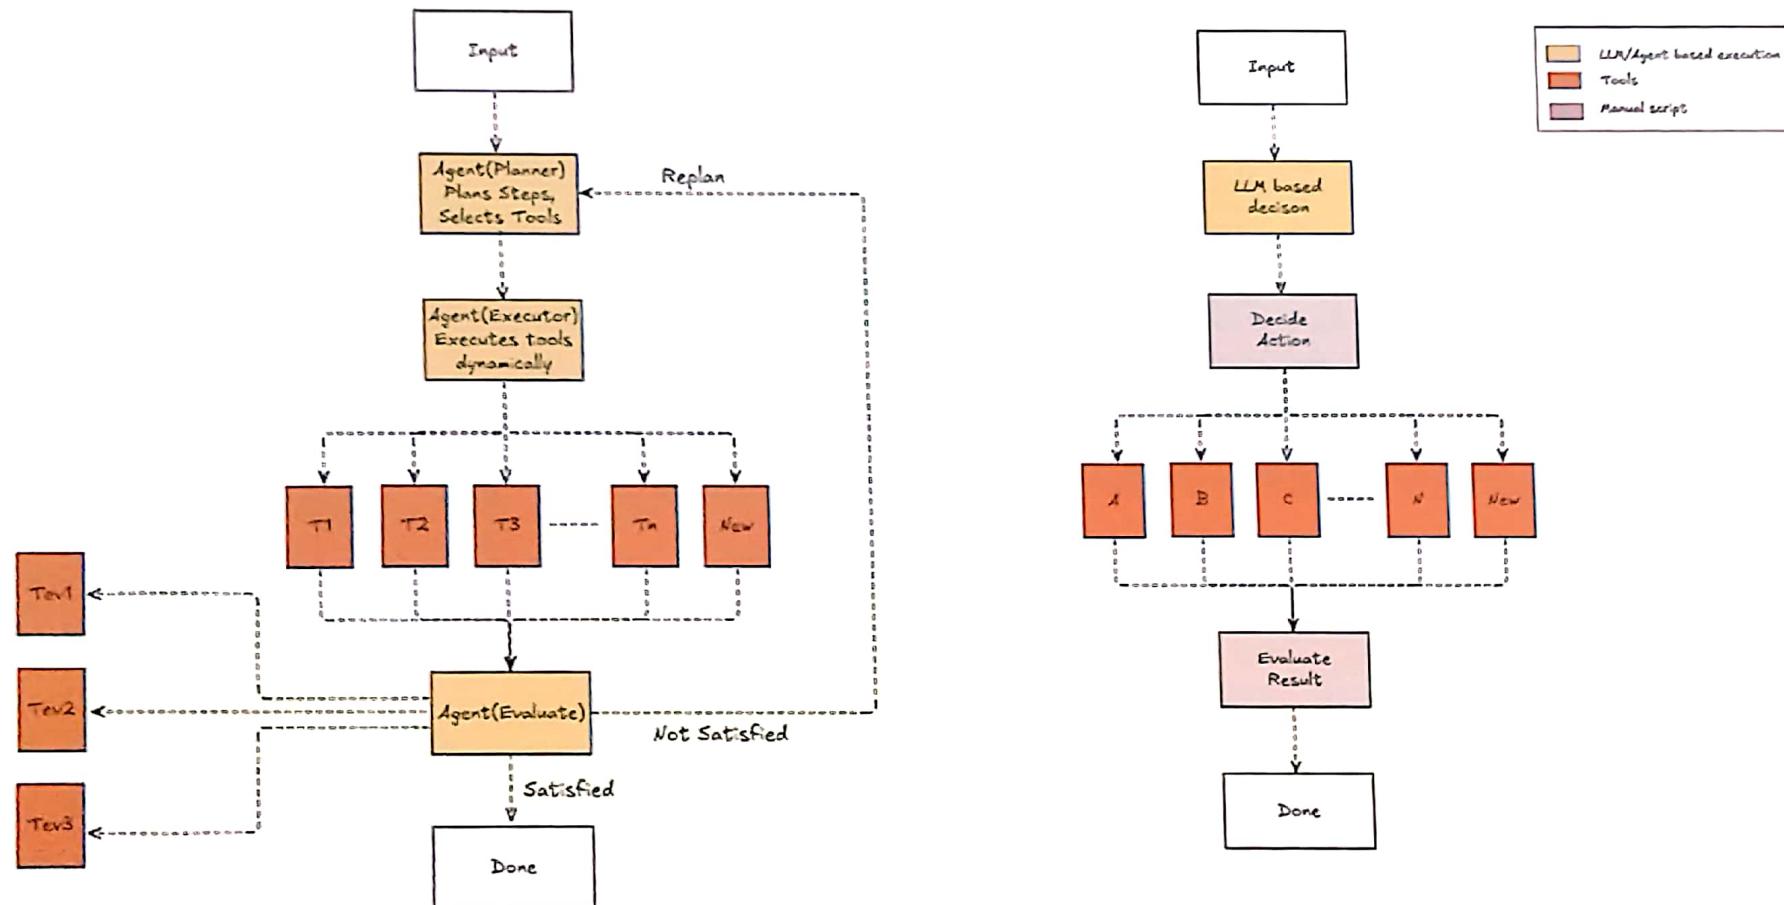


# Agents vs Workflows: A Turing Test for Your Architecture

Samir Kumar Nayak



# Agents vs LLM based fixed workflows



# Problem: Should i use an AI Agent or Simple Pipeline?

When building AI powered applications, we face this critical decision.

Wrong choice = wasted time, money or poor results!

Why not ask ChatGPT?

ChatGPT says:

- “It depends”...
- Gives an opinion which might change each time
- Black box reasoning, no metrics how the opinion is formed



# The Turing Test Framework: PRISM

## PRISM: Problem Recognition and Intelligent Strategy Mapping

A framework that analyzes workflow problems and recommends whether to implement your LLM based workflow as a:

- Fixed Workflow
- Agentic Workflow
- Hybrid Approach

The framework models your problem as a graph of nodes (processing steps) and edges (transitions), then analyzes structural properties to provide a recommendation with confidence scores.



PRISM  
Problem Recognition and Intelligent Strategy Mapping  
Analyze workflow problems to determine optimal implementations:  
Fixed workflow • Agentic Workflow • Hybrid Approach

```
usage: prism.py [-h] [-t] [-template] [-interactive] [-examples] [-output-dir OUTPUT_DIR] [--no-radar]
                 [-version] [-generate DESCRIPTION] [-domain DOMAIN] [-no-analyze] [-react] [-verbose] [-max-steps MAX_STEPS]
                 [-research] [-research-depth {quick,standard,deep}]
                 [yaml_file]

Problem Recognition and Intelligent Strategy Mapping

positional arguments:
  yaml_file      Path to YAML problem definition file

options:
  -h, --help       show this help message and exit
  -t, --template   Print YAML template to stdout
  -i, --interactive   Run interactive problem definition wizard
  -e, --examples    Analyze built-in example problems
  -o, --output-dir OUTPUT_DIR, --OUTPUT_DIR
                    Output directory for generated files
  --no-radar        Skip radar chart generation
  -v, --version     Show version information
  -g, --generate DESCRIPTION, --g DESCRIPTION
                    Generate YAML from natural language description (requires NVIDIA_API_KEY)
  -d, --domain DOMAIN, --D DOMAIN
                    Domain hint for generation (e.g., e-commerce, customer_support)
  -n, --no-analyze   Skip automatic analysis after generation
  -r, --react        Use ReAct agent approach for generation (more thorough)
  -v, --verbose      Show reasoning process during generation
  -m, --max-steps MAX_STEPS
                    Maximum ReAct steps for --react mode (default: 10)
  -s, --research     Use deep research to enhance generation (searches web for best practices)
  -d, --research-depth {quick,standard,deep}
                    Depth of research: quick (2 queries), standard (4), deep (6)

Examples:
# Analyze a YAML file
prism my_problem.yaml

# Generate YAML from natural language (requires NVIDIA_API_KEY)
prism -g "Build a customer support chatbot"

# Generate with deep research (requires TAVILY_API_KEY)
prism -g "Build an order processor" --research --verbose

# Research with custom depth (quick/standard深深)
prism -g "Build a fraud detection system" --research --research-depth deep

# Combine research with ReAct agent for best results
prism -g "Build an AI assistant" --research --react --verbose

# Generate template
prism -t > my_problem.yaml

# Run interactive mode
prism -i

# Analyze example problems
prism -e
```



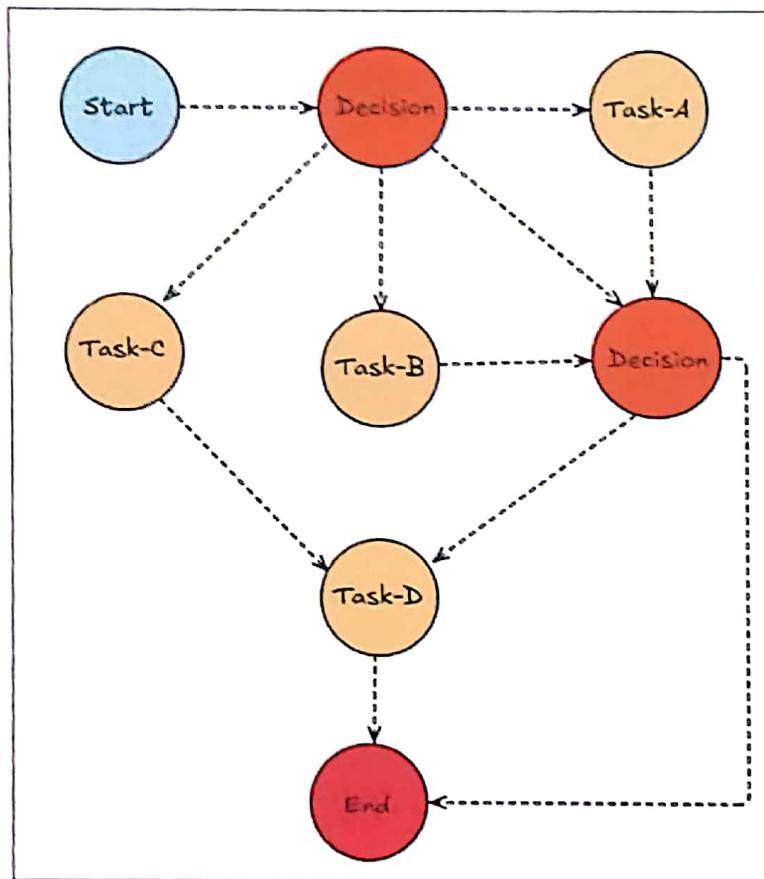
# The Problem-First Approach

Before writing a single line of code:

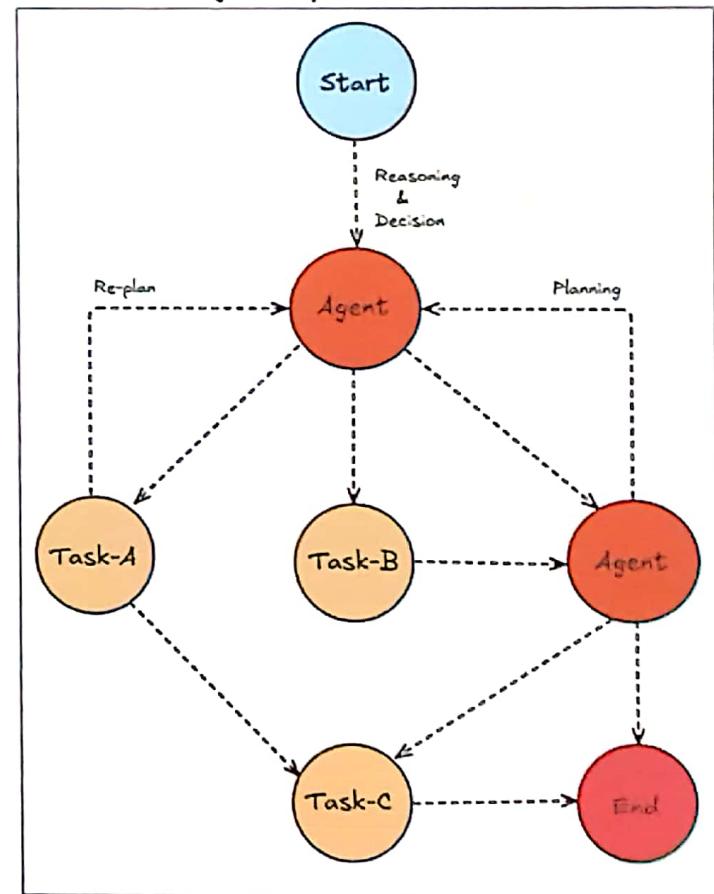
- Map the problem as a graph
- Problem = Series of States(Nodes) + Transitions(Edges)
- Architecture is defined by how we move between nodes
- The right choice depends on: Who decides the transition?

# Fixed & Agent based workflow (Problem breakdown)

Fixed LLM based workflow

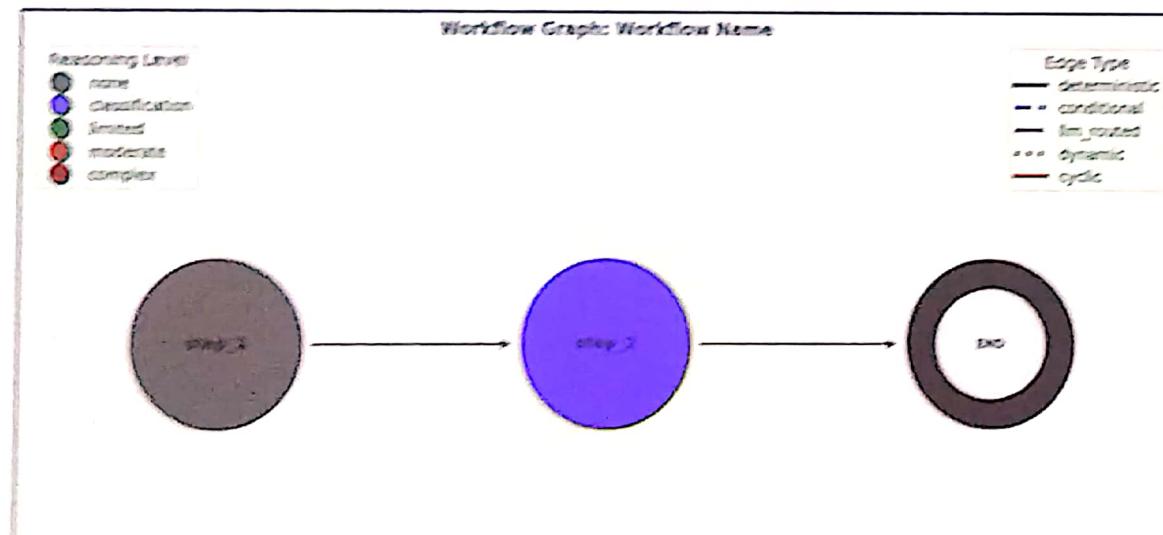


Agent/Hybrid based workflow



# Structure the problem statement

Section	Purpose
Problem	Name, description, and domain classification
Constraints	Latency limits, determinism, auditability requirements
Steps	Processing nodes with input/output formats and reasoning levels
Edges	Connections between steps with transition types

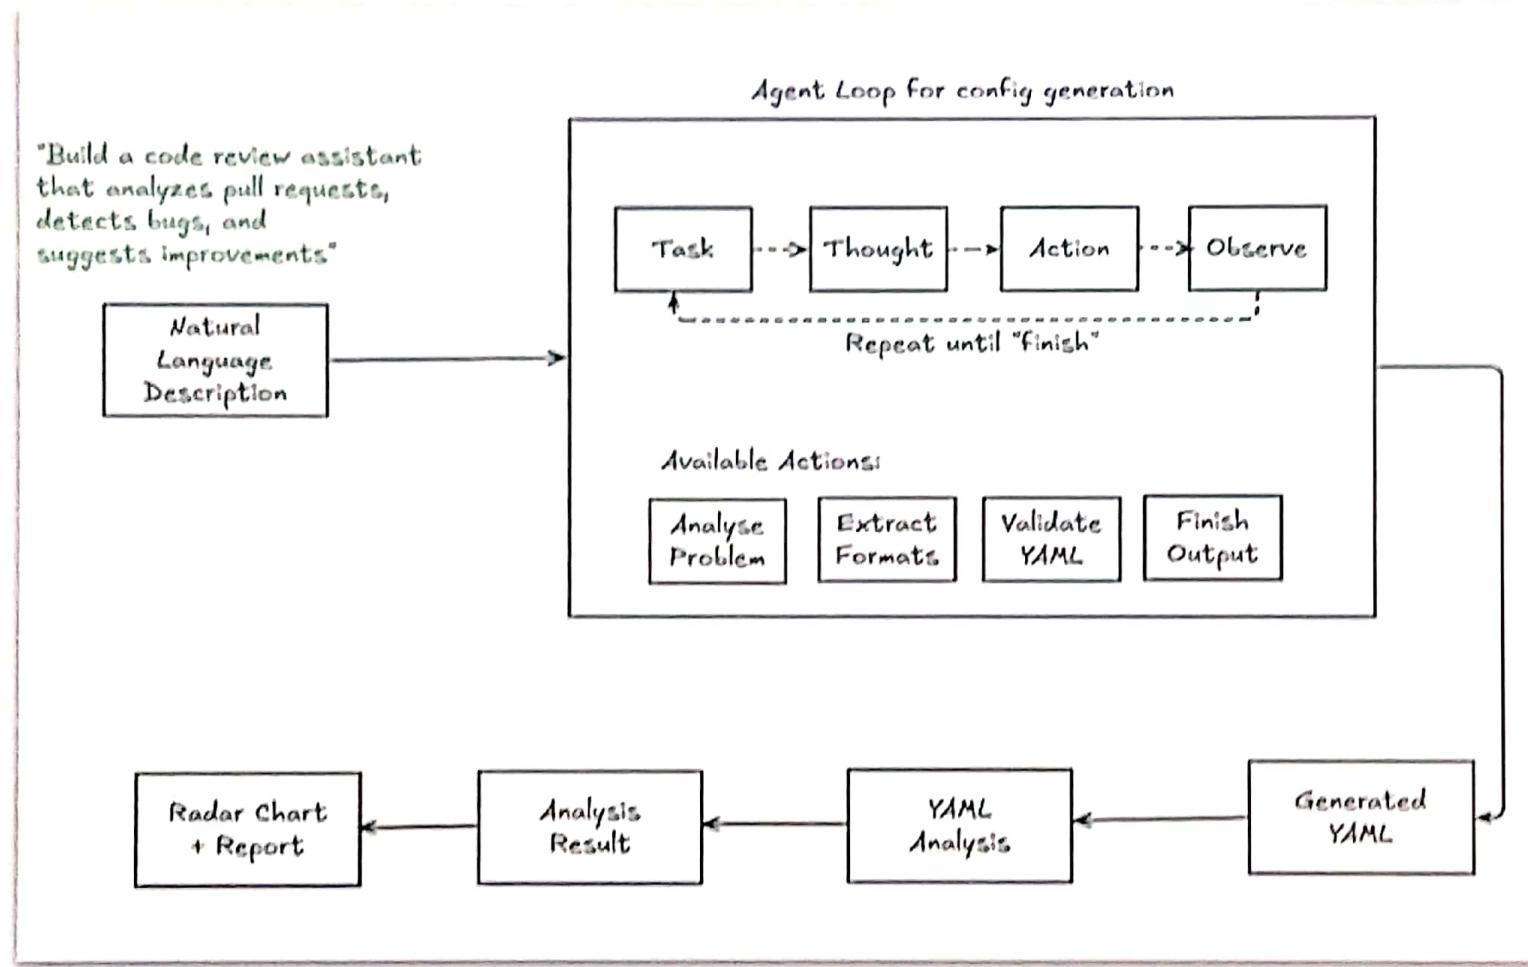


```

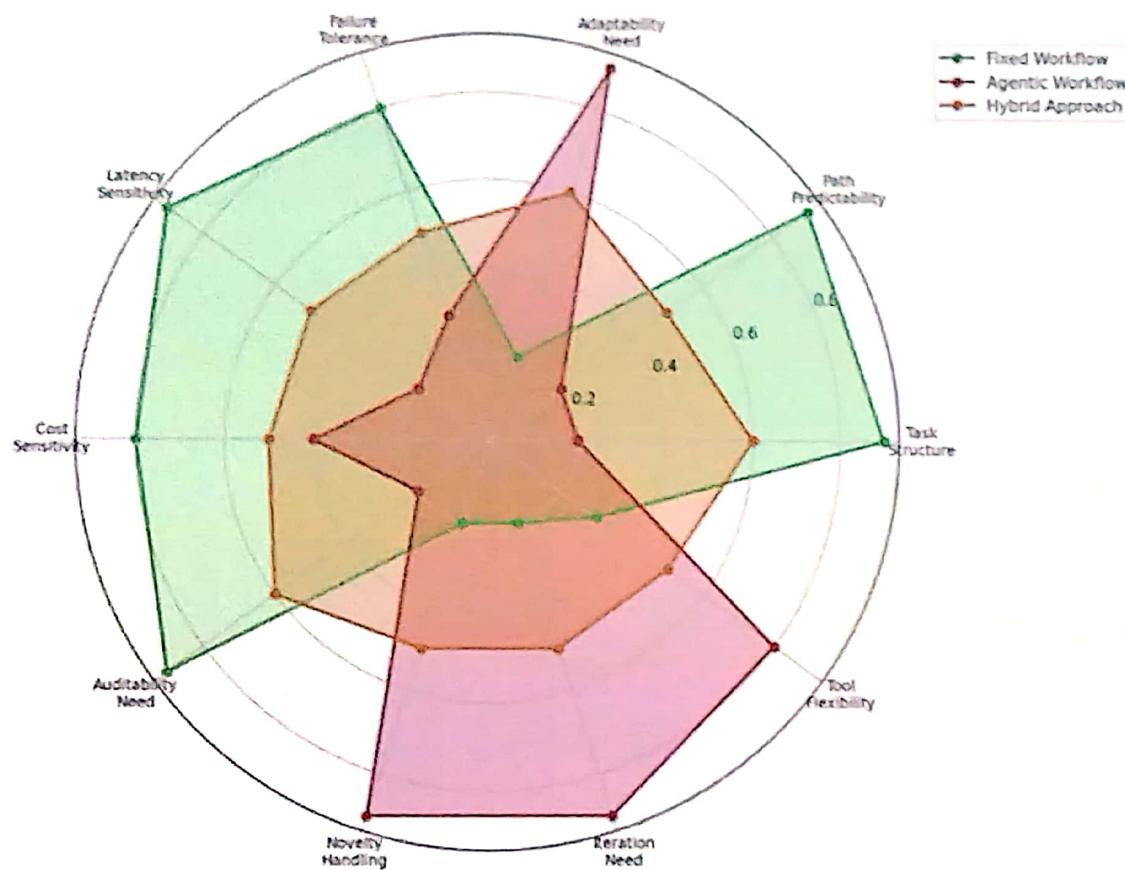
# PRISM Workflow Template
#
# problem
#   name: "My Workflow"
#   description: "What this workflow does"
#
# constraints
#   latency_ms: null      # null | <milliseconds>
#   requires_determinism: false # true | false
#   requires_auditability: false # true | false
#
# steps
#   - name: "step_1"
#     description: "First step"
#     input_formats: "structured" # structured | semi | unstructured
#     output_formats: "structured" # structured | semi | unstructured
#     reasoning: "none" # none | classification | limited | moderate | complex
#     requires_retry: false # true | false
#     requires_self_eval: false # true | false
#
#   - name: "step_2"
#     description: "Second step"
#     input_formats: "structured"
#     output_formats: "structured"
#     reasoning: "classification"
#     requires_retry: false
#     requires_self_eval: false
#
# edges
#   - from: "step_1"
#     to: "step_2"
#     type: "deterministic"
#     conditions: []
#     max_iterations: null
#
#   - from: "step_2"
#     to: "END"
#     type: "deterministic"
  
```



# PRISM Workflow



# Workflow dimension profiles



#	Dimension	What It Measures	Favors
1	Adaptability Need	Need to handle changing inputs	Agentic
2	Path Predictability	How deterministic the execution path is	Fixed
3	Task Structure	How well-defined the requirements are	Both
4	Tool Flexibility	Benefit from iterative tool selection	Agentic
5	Iteration Need	Need for loops and retries	Agentic
6	Novelty Handling	Handle unexpected/edge-case scenarios	Agentic
7	Auditability Need	Tracing and compliance requirements	Fixed
8	Cost Sensitivity	Importance of minimizing API costs	Fixed
9	Latency Sensitivity	Importance of fast response time	Fixed
10	Failure Tolerance	Need for graceful error handling	Fixed

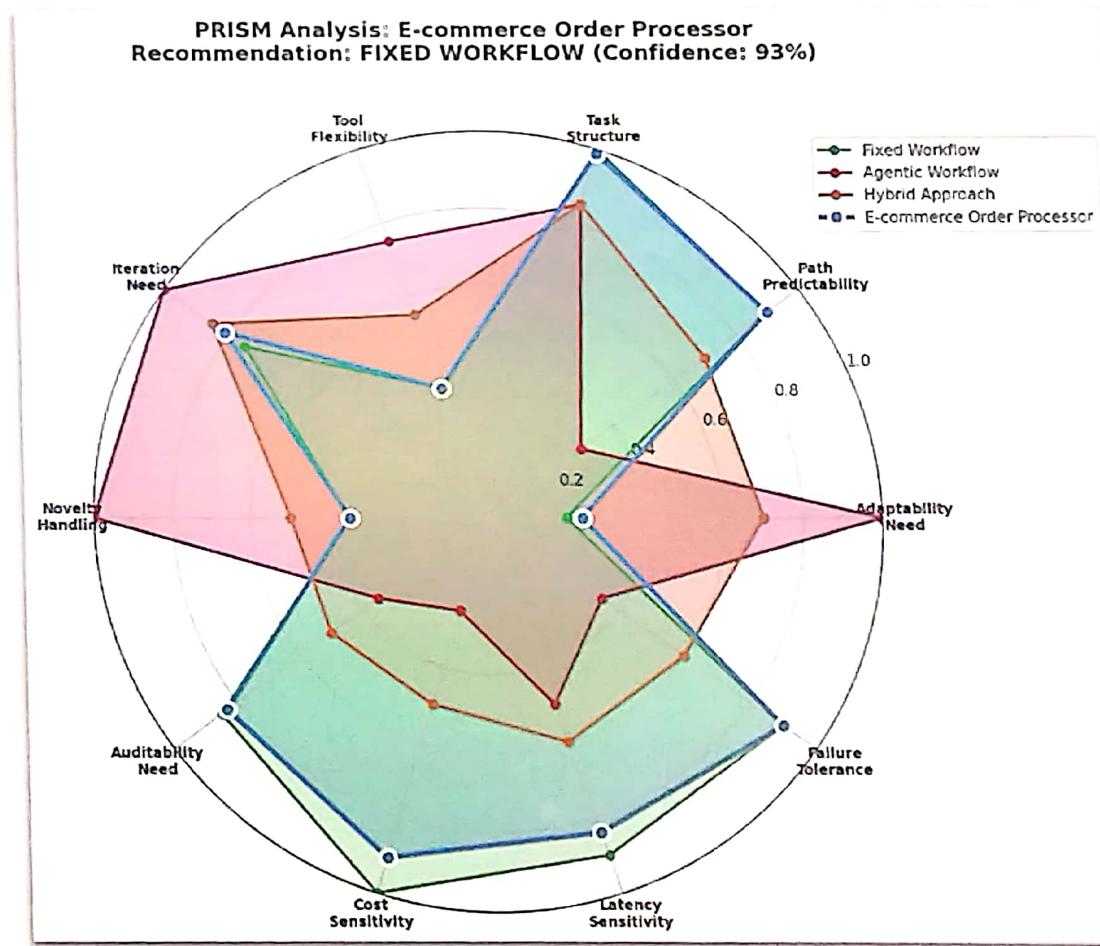
# Examples: e-commerce order processing

```
# PRISM: E-commerce Order Processing Workflow
problem
  name "E-commerce Order Processor"
  description "Processes incoming orders from validation to fulfillment"
  domain "e-commerce"

constraints
  latency_ms: 2000
  requires_determinism: true
  requires_auditability: true

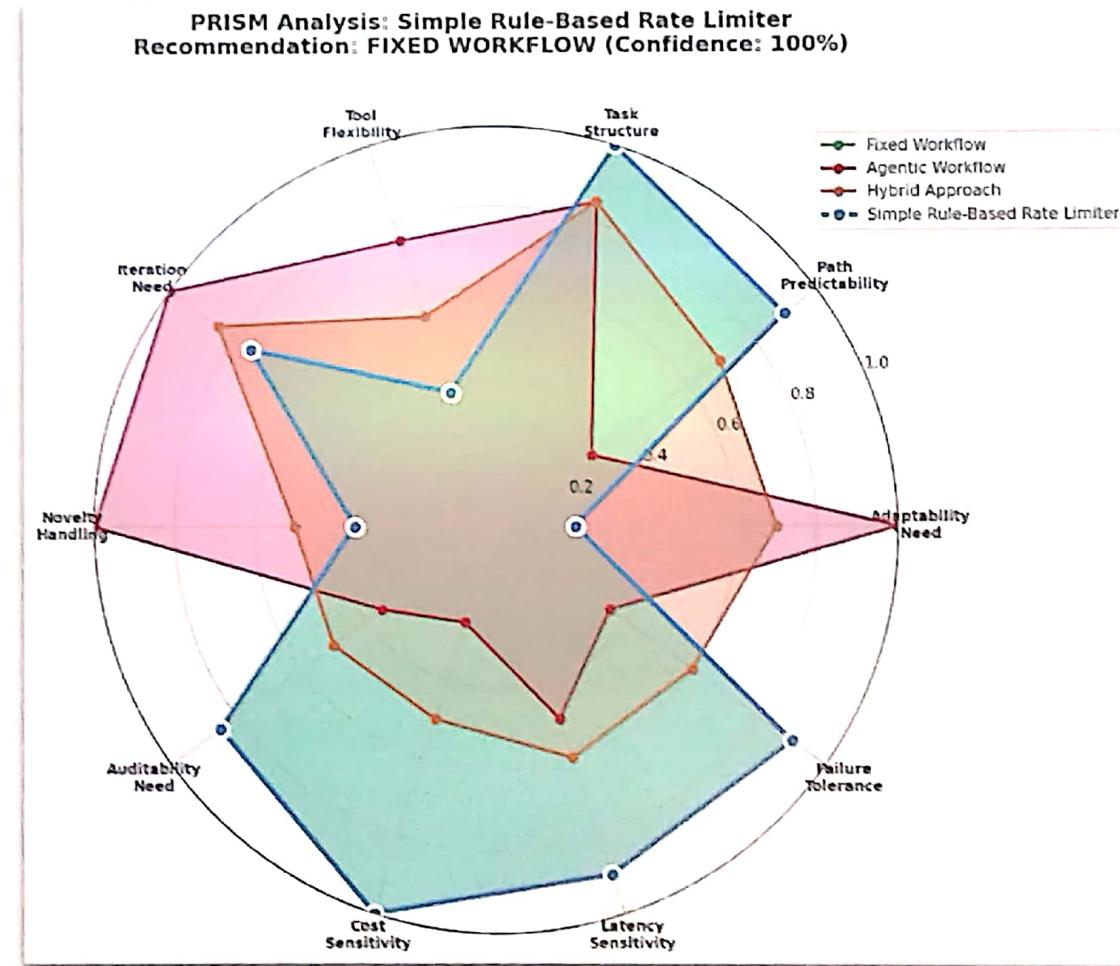
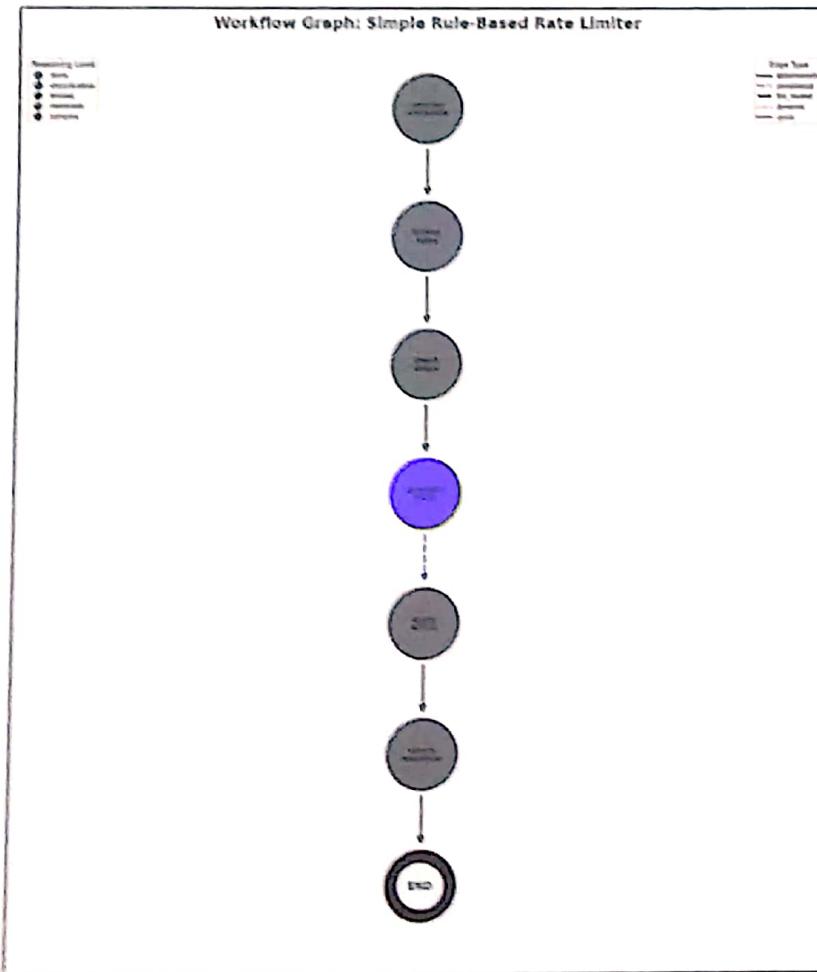
steps
  - {name: validate_order, input_format: semi, output_format: structured, reasoning: classification,
    requires_retry: false, requires_self_eval: false}
  - {name: check_inventory, input_format: structured, output_format: structured, reasoning: none,
    requires_retry: false, requires_self_eval: false}
  - {name: calculate_pricing, input_format: structured, output_format: structured, reasoning: limited,
    requires_retry: false, requires_self_eval: false}
  - {name: process_payment, input_format: structured, output_format: structured, reasoning: none,
    requires_retry: true, requires_self_eval: false}
  - {name: handle_payment_failure, input_format: structured, output_format: structured, reasoning: limited,
    requires_retry: true, requires_self_eval: false}
  - {name: create_shipment, input_format: structured, output_format: structured, reasoning: none,
    requires_retry: false, requires_self_eval: false}
  - {name: send_confirmation, input_format: structured, output_format: structured, reasoning: limited,
    requires_retry: false, requires_self_eval: false}

edges
  - {from: validate_order, to: check_inventory, type: conditional, conditions: [order_valid=true]}
  - {from: validate_order, to: END, type: conditional, conditions: [order_valid=false]}
  - {from: check_inventory, to: calculate_pricing, type: deterministic}
  - {from: calculate_pricing, to: process_payment, type: deterministic}
  - {from: process_payment, to: create_shipment, type: conditional, conditions: [payment_success=true]}
  - {from: process_payment, to: handle_payment_failure, type: conditional, conditions:
    [payment_success=false]}
  - {from: handle_payment_failure, to: process_payment, type: cyclic, max_iterations: 3}
  - {from: handle_payment_failure, to: END, type: conditional, conditions: [max_retries_exceeded]}
  - {from: create_shipment, to: send_confirmation, type: deterministic}
  - {from: send_confirmation, to: END, type: deterministic}
```

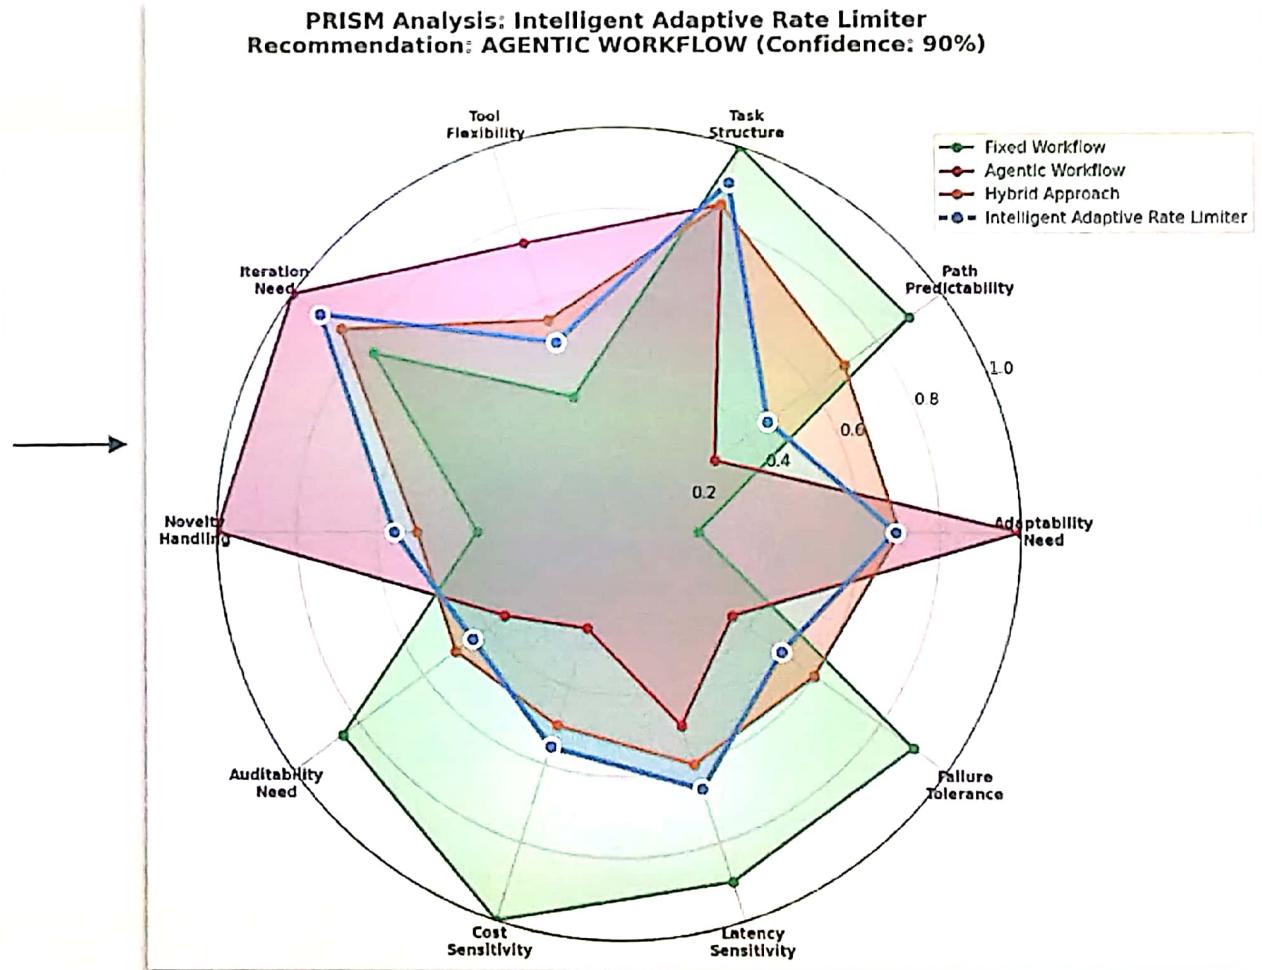
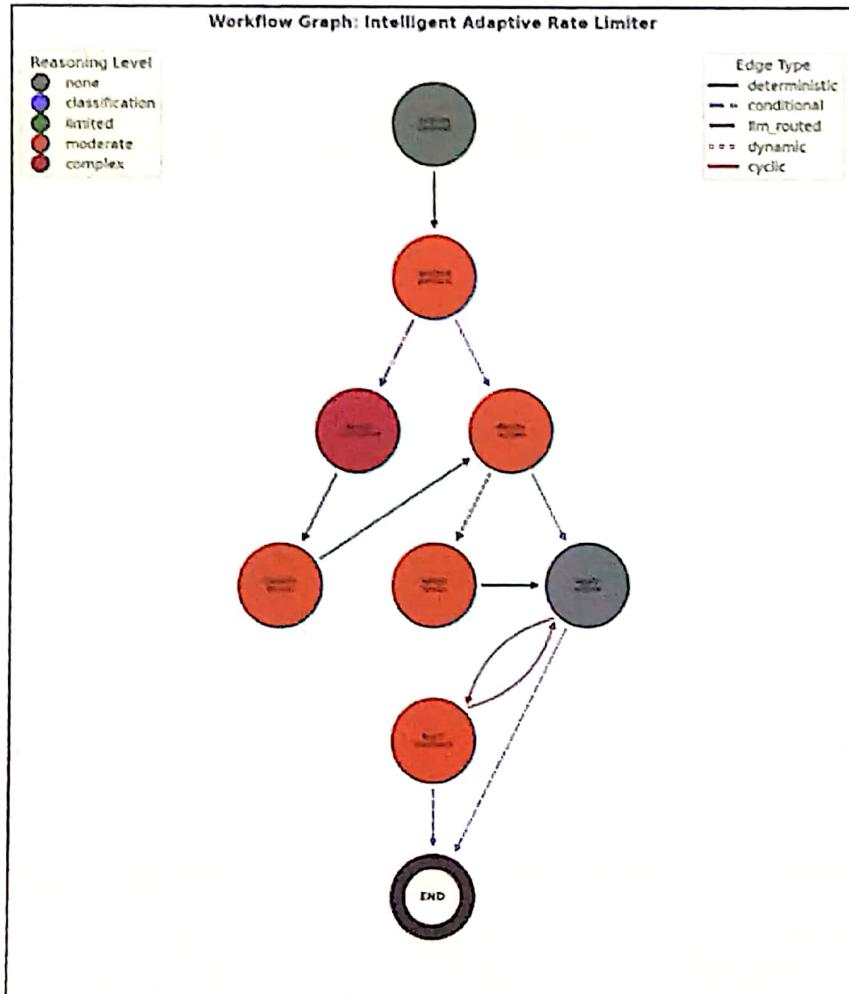




# Examples: Simple rate limiter for custom API



# Examples: Intelligent rate limiter for custom API



# Dimension Score Calculation from YAML

Formula: score = base + factor × (weighted sum of inputs)

Dimension	Favors	Base	Factor	Range	Weight Components
Adaptability Need	Agentic	0.2	0.8	0.2–1.0	inverse_data_format: 0.35, inverse_reasoning: 0.35, dynamic_edges: 0.15, self_eval_nodes: 0.15
Path Predictability	Fixed	0.3	0.6	0.3–0.9	paths_enumerable: 0.4, inverse_dynamic_edges: 0.3, no_unbounded_cycles: 0.3
Task Structure	Both	0.85	0.15	0.85–1.0	data_format: 0.5, reasoning: 0.3, inverse_dynamic_edges: 0.2
Tool Flexibility	Agentic	0.35	0.4	0.35–0.75	llm_routed_edges: 0.3, dynamic_edges: 0.3, inverse_reasoning: 0.2, complex_nodes: 0.2
Iteration Need	Agentic	0.75	0.25	0.75–1.0	has_cycles: 0.25, unbounded_cycles: 0.25, retry_nodes: 0.25, self_eval_nodes: 0.25
Novelty Handling	Agentic	0.35	0.65	0.35–1.0	unstructured_ratio: 0.35, inverse_reasoning: 0.35, exploration: 0.15, dynamic_edges: 0.15
Auditability Need	Fixed	0.35	0.5	0.35–0.85	paths_enumerable: 0.35, data_format: 0.25, topology: 0.25, inverse_dynamic_edges: 0.15
Cost Sensitivity	Fixed	0.25	0.75	0.25–1.0	reasoning: 0.35, no_cycles: 0.25, inverse_complex_nodes: 0.25, inverse_retry_nodes: 0.15
Latency Sensitivity	Fixed	0.5	0.4	0.5–0.9	reasoning: 0.3, no_cycles: 0.3, inverse_dynamic_edges: 0.2, topology: 0.2
Failure Tolerance	Fixed	0.35	0.55	0.35–0.9	paths_enumerable: 0.3, no_unbounded_cycles: 0.25, no_complex_reasoning: 0.25, topology: 0.2

All weights are heuristically designed, not empirically validated([weights](#))



Scanned with OKEN Scanner

# Intermediate Metrics

Metric	Source YAML Fields	Formula
data_score	steps[].input_format, steps[].output_format	Count of steps where both are "structured" ÷ total steps
reasoning_score	steps[].reasoning	Count of steps with "none", "classification", or "limited" ÷ total steps
dynamic_edges	edges[].type == "dynamic"	Count
llm_routed_edges	edges[].type == "llm_routed"	Count
has_cycles	edges[].type == "cyclic"	true if any exist
unbounded_cycles	edges[].type == "cyclic" + edges[].max_iterations == null	true if any cyclic edge has no limit
paths_enumerable	edges[], edges[].max_iterations	true if no dynamic edges AND no unbounded cycles
complex_nodes	steps[].reasoning in ["moderate", "complex"]	Count
retry_nodes	steps[].requires_retry == true	Count
self_eval_nodes	steps[].requires_self_eval == true	Count
unstructured_ratio	steps[].input_format, steps[].output_format	(unstructured inputs + outputs) ÷ (total_nodes × 2)

# PRISM: Key Takeaways

The problem we are trying to solve:

**"Should I build this as a fixed workflow or let an agent figure it out?"**

This question costs teams weeks of trial-and-error, over-engineered solutions, and missed performance targets.

The Core Insight:

- Not every LLM workflow needs to be agentic.
- Not every workflow can be fixed.
- The problem structure determines the right approach.

PRISM Approach:

- Model: Represent any workflow as a structured graph of steps and transitions
- Measure: Score in 10 workflow dimension metrics that differentiate Fixed from Agentic
- Decide: Get a quantified recommendation, not a guess

\*The dimension weights are heuristically designed based on domain expertise and logical relationships between workflow characteristics. They follow consistent patterns (primary vs secondary factors) and prevent edge-case scores. Future work could empirically validate these weights against real-world implementation outcomes.